

Swipe away junk

gm legends. It's Sunday.

This week: swipe right on all the newest email clients, where Meta's latest LLM ranks against the competition, how to pitch a VC so they'll actually listen, and where to send your agents when you don't have a terminal open. Plus, five of our favorite launches from the past week.

You don't need to swipe on this email, legend. Just scroll down. Enjoy.


P.S. Launching soon? We'd love to hear about it →

editorial@producthunt.co



IN THE NEWS

Tinder for email



Tinder's great, isn't it?

You swipe left, and you never see the person again. Or you swipe right...and you never see the person again (probably). Amazing!

A new iOS mobile app, **Avec**, brings the swiping system to email—but only gets rid of the emails you're done with. A left swipe means you need to handle it; a right swipe archives it. You can then go through and voice-dictate replies.

Avec isn't the first tool to reimagine email. Think: **Superhuman**, a productivity-centric email client that promises to save people 4 hours a week. Or **Shortwave**, a fast, AI-driven client with tons of keyboard shortcuts.

In fact, Product Hunt has seen several client launches in recent weeks, including:

- **AgentMail**, an inbox for your AI agent
- **Apparent for Gmail**, a free add-on to Gmail that gives it all the features you wish it had
- **BAREMAIL**, a stripped-down Gmail client for when you have crappy WiFi
- **InboxAgents**, which fixes the notoriously prickly LinkedIn inbox, then unifies it with your email and messages from other platforms
- **Inbox Autopilot** by Dimension, which peruses your inbox like an office assistant and drafts responses for you
- **ReplylessAI**, a budget-friendly app for hitting "inbox zero"

Find your client

FROM THE FORUMS

20 tips for pitching to VCs



Not so long ago, **Dan Bulteel** was an exec at Adidas and then TikTok, where he ran social media marketing globally. Dude knows his stuff. He's mentored early-stage founders for Techstars, and he advises multiple companies on their marketing. But before he built **Meet-Ting**, an AI-powered app for managing your calendar, he'd never really been through the funding gauntlet as a founder.

After pitching London and San Francisco VCs, he has some notes.

Here's everything he'd tell a founder before their first VC call, from the big-picture stuff (why you should act like the belle of the ball) to the small hacks others might overlook (why to use DocSend).

[Get the tips](#)

Ad

Are you really still typing?

Full disclosure: Wispr Flow is the AI dictation tool most of us at Product Hunt (use we still have a few holdout typers, what romantics). Hold a key, talk, and clean text drops straight into whatever app you're already in — Slack, email, Notion, your IDE, wherever your cursor lives. No switching windows. No copy-paste ritual. Just say the thing – yes, you can whisper it – and even your most run-on sentences will be turned into polished writing at 4x the speed of typing.

Type 4x Faster

89% of messages sent
without a single edit.




 **Flow** | Voice-to-text that actually gets it right.

Try Flow for free

WHY I BUILT THIS

Give your agent space



Thursday's **Alpha Day** brought tons of truly fresh products that haven't been seen elsewhere. One of these was **Grass**, built by serial maker Sunny. He writes:

"...we noticed something weird. We message our teammates from anywhere, on the train, at lunch, from the couch. But the moment we want to talk to our coding agent, we need a laptop and a terminal open. That felt broken. Coding agents aren't tools you operate anymore. They're colleagues you work with."

Keep reading to see his solution.

Touch Grass

IN THE NEWS

Meta's got a new model



Meta this week announced its latest large language model, **Muse Spark**. The LLM, which select users can access through a private API preview, fills the shoes left vacant by **Llama 4**. Meta moved on from Llama 4 after disappointing uptake.

The biggest difference?

Whereas Llama 4 was open-source, Muse Spark is closed, putting it on a collision course with Google **Gemini**, Anthropic's Claude **Opus**, OpenAI's **ChatGPT**, and xAI's **Grok**. According to **benchmark** test results released by Meta, Muse Spark is "competitive" with all four but lands in the middle overall.

Meta has gone to great lengths to win the AI race, bringing Alexandr Wang over from Scale AI, aggressively headhunting engineers, and, recently, shifting its focus away from the metaverse.


Now, it'll need to do a bit more to pull people away from their preferred models. While it's prepping its launch, competitors aren't standing pat.

This week alone, OpenAI launched **ChatGPT on CarPlay**, and Anthropic came out with **Claude Managed Agents** and its **Claude Advisor** tool.

[Compare LLMs](#)

Weekly

Leaderboard highlights



Brila — One-page websites from real Google Maps reviews



1206

Brila builds one-page websites from Google Maps reviews. Paste in a business, and it pulls out the patterns behind why people actually pick that place, then turns that into copy, structure, and imagery for a site that sounds like real customers instead of a guy typing trusted local service into ChatGPT. It is doing the content part first, which is already a better sign than half this category.



Velo — Share anything as video messages



666

Velo takes your messy screen recordings, links, or slides and turns them into something watchable. You can record, upload, or just paste a URL, and it builds a cleaner version with a script, voiceover, and visuals that actually line up.



Offsite — Build teams of humans and agents, watch them work.



575

Offsite puts humans and AI agents on the same team instead of scattering them across tabs, terminals, and whatever other mess you have open. You drop everyone into a live org chart, connect them, and watch the work move between people and agents in real time. It plugs into tools like Claude Code, OpenClaw, and other MCP-compatible agents, with approvals in the loop so nobody quietly goes rogue.



Wispr Flow: Dictation That Works Everywhere — Stop typing. Start speaking. 4x faster.



Promoted  527 •  Productivity

2129



SuperShrimp — Fix your terrible posture



SuperShrimp uses your MacBook camera to catch you slouching while you work. It gives you a live posture score, sends alerts when you start folding into your laptop, keeps everything local and offline, and adds XP so better posture actually levels something up.

332



riffle — An infinite, collaborative playground for music creation

Riffle is a browser-based music playground for making songs with other people without getting buried in pro-audio nonsense. You can build ideas with samples, instruments, and audio, bring friends into the same space, and use its AI sous-chef for help with things like critique, guidance, and shaping a track without handing the whole job over.

Hi there,

Productivity looks different when you can write at the speed of thought. Meet [Monologue](#) (available for MacOS and iOS), our AI voice dictation tool. It helps you work 3x faster using perfect voice dictation that knows what you mean in every situation.



Here are three Monologue features that make it so powerful:

- **Auto formatting:** It reformats its output based on what and where you're writing.
- **Deep context:** It can “see” your screen and understand what you're referencing. When you say “Current Step Index” out loud, Monologue writes “currentStepIndex.”
- **Auto dictionary:** It learns your vocabulary and adds words to its dictionary automatically, so you can speak naturally and stay in the flow.

Try Monologue

The shortest distance between two points is a straight line. The shortest distance between your thoughts and the page is Monologue.

If you've been around AI for any length of time, you've come across the idea of prompts (it's how you give instructions to an LLM like ChatGPT).

As it turns out, there are things you can do to better craft your prompts to improve the quality of results. That craft is known as "prompt engineering"--and it's exactly what it sounds like, engineering a prompt so the LLM does what you want.

But here's what's fascinating: as context windows have exploded from 4K to over 1M tokens in the last two years, something more powerful than prompt engineering has emerged: *Context Engineering*. (h/t to my friend Tobi from Shopify for coining this, I think)

The shift may seem simple but the framing is completely different--instead of optimizing how you ask, you're now optimizing what the AI has access to when it *thinks*.

So in today's newsletter, I'm writing a more "instructional" post on:

- - How context windows work and why they're important
- The evolution of prompt engineering to context engineering
- Why this shift matters for anyone building with AI


—[@dharmesh](#)



What is a Context Window?

Think of it as a big sheet of paper that you pass to the LLM.

You can write a certain number of words on it. Those words can be "write me a limerick about pickleball." Or, they can be "Here's a 5,000-word essay, give me a 100-word summary" (and you include the whole essay on that big sheet of paper).



I thought this was funny (thanks GPT-4o)

Context windows have a limit, usually expressed as a number of "tokens." A token is roughly $\frac{3}{4}$ of an English word or about four characters. "ChatGPT" is two tokens: "Chat" and "GPT." This matters because billing, latency, and memory all scale with token count.

Here's a key limitation with LLMs: they only know what they were trained on and what's provided to them in the context window.

But as context windows have exploded in token size over the last 2 years, we've gotten increasingly clever about what we put in that context window, such as:

How AI "Remembers" Your Conversation

You may have noticed that ChatGPT and Claude have good short-term memory. That's why you can ask follow-up questions and they know what you're talking about. The way that happens? Behind the scenes, they're actually passing your prior prompts and the outputs into the context window. It's kind of like that movie *Memento*, where the main character forgets everything, and has to write notes on his body so he can remember who he is and what's happening (great movie by the way).

RAG: Teaching AI on Demand

There's an approach called RAG (Retrieval-Augmented Generation) where we find a set of documents that are related to the user prompt, and then pass those into the context window (thereby "teaching" the LLM precisely the things that will help it answer a question or prompt).

Tool Calling: Extending AI's Capabilities

We've also introduced the idea of tool calling. It's a way to tell the LLM: "Hey, pretend like you have access to these set of tools for things like web search, stock price lookups, weather forecasts...whatever." LLMs can then use one or more of the tools provided if it will help them process a prompt.

Here's a quirky but delightful detail on how tool calling works: When we let the LLM know what tools it has available, it doesn't actually *call* the tools itself. The core LLM isn't built to do that. It just takes context in and produces output. We tell it what tools it has access to, and the output, it lets the application (like ChatGPT) which tool it would like to invoke. The app then invokes the tool on the LLMs behalf and puts the results in the context window and sends it back. It's a clever way to work around the fact that LLMs just work with a context window.

With a million-token context window, you can feed an AI an entire codebase, a complete business plan, or months of customer support conversations, and it can reason across all of that information simultaneously.



Why Context Engineering Matters

Remember the prompt engineering craze of 2023? Headlines were pushing \$300k+ salaries for "Prompt Engineers" at places like Anthropic.

A lot of the jobs went viral because they required no formal CS degree -- just the ability to speak fluent "LLM."

**Prompt Engineer and Librarian**
Anthropic - San Francisco, CA
\$175,000 - \$335,000 a year

Anthropic's mission is to create reliable, interpretable, and steerable AI systems. We want AI to be safe for our customers and for society as a whole.

Anthropic's AI technology is amongst the most capable and safe in the world. However, large language models are a new type of intelligence, and the art of instructing them in a way that delivers the best results is still in its infancy — it's a hybrid between programming, instructing, and teaching. We want you to figure out the best methods of prompting our AI to accomplish a wide range of tasks, then document these methods to build up a library of prompts and a set of tutorials that allows others to learn prompt engineering or simply find prompts that would be ideal for them.

Given that the field of prompt-engineering is arguably less than 2 years old, this position is a bit hard to hire for! If you have existing projects that demonstrate prompt engineering on LLMs or image generation models, we'd love to see them. If you haven't done much in the way of prompt engineering yet, you can best demonstrate your prompt engineering skills by spending some time experimenting with Claude or GPT3 and show that you've managed to get complex behaviors from a series of well crafted prompts.

Representative Projects

- Discover, test, and document best practices for a wide range of tasks relevant to our customers.
- Build up a library of high quality prompts or prompt chains to accomplish a variety of tasks, with an easy guide to help users search for the one that meets their needs.
- Build a set of tutorials and interactive tools that teach the art of prompt engineering to our customers.
- Work with large enterprise customers on their prompting strategies.

Old screenshot of Anthropic's Prompt Engineer job listing

While some prompt engineers were hired, it didn't end up being the massive job category many predicted. The reason? Everyone became a prompt engineer.


And now, there's something even more important emerging: Context Engineering.

Context Engineering is essentially a higher-level version of Prompt Engineering.

An (overly) simple explanation: Prompt engineering was like learning to ask really good questions. Context engineering is like being a librarian who decides what books someone has access to before they even start reading.

What Context Engineers Actually Do:

- - **Curate:** Decide which documents, memories, or APIs matter for each specific task
- **Structure:** Layer system messages → tools → retrieved data → user prompt in optimal order
- **Compress:** Summarize or chunk information to stay under token limits while preserving what matters
- **Evaluate:** Measure accuracy and watch for "context dilution" where irrelevant info distracts the model



Keep in mind, more context means richer documents and longer conversations. But cost and latency rise roughly linearly with window length. This leaves a ton of room for Context Engineers to discover best practices (many of which are still emerging).

Context engineering requires thinking about information architecture, data strategy, and user experience in ways prompt engineering never did. The companies who master this will have a massive competitive advantage. As I like to say: Prompting tells the model how to think, but context engineering gives the model the training and tools to get the job done.



The Context-First Future

The prompt engineering hype of 2023 taught us an important lesson: the most valuable AI skills aren't about learning secret phrases or clever tricks.

They're about understanding how to architect intelligent systems that have access to the right information at the right time.


It's a fundamental shift from optimizing *sentences* to optimizing *knowledge*. This matters for anyone building with AI today.

Personally, as someone building in this space with [agent.ai](#), I'm constantly amazed by what becomes possible when you stop thinking about AI as a chatbot and start thinking about it as a reasoning engine with access to the right context and tools.

The prompt engineering era taught us to talk to AI. The context engineering era is teaching us to think with AI.

And if you made it this far, I hope you found this "instructional-style" newsletter helpful. It's a bit different from my usual format, I'd love to know what you think in the poll below.

—Dharmesh ([@dharmesh](#))



Hello, and happy Sunday! Two housekeeping notes: Our next cohort of [Claude Code for Absolute Beginners](#) is taking place on Tuesday, April 14, and Every has opened [seven new roles](#). Join us!—[Kate Lee](#)

Was this newsletter forwarded to you? [Sign up](#) to get it in your inbox.

Knowledge base

[“Writing With AI Is Harder Than You Think”](#) by *Katie Parrott/Working Overtime*: The discourse about AI and writing generally assumes prompt in, text out, done. **[Katie Parrott](#)** shows her much more involved process: an agent that interviews her before she writes a word, a back-and-forth on her structure that she has to fight for, a panel of AI critics named Hemingway and Hitchcock, and a last read that flags anything that sounds machine-generated. Read this because successful AI writing demands more judgment, not less.

[“Your Best AI Strategy Starts at the Top”](#) by *Natalia Quintero and Mike Taylor*: Most executives approach AI like a software purchase—evaluate, compare features, and plug in.

Natalia Quintero and **Mike Taylor** see it differently: Using AI is people management, not platform adoption. You delegate clearly, check the output, and supply the judgment the model doesn't have. Read this for the five concrete actions senior leaders can take to increase AI adoption within their companies.

“Get Your Hands Dirty” by *Every Staff/Context Window*: Anthropic blocked Claude subscriptions from working with third-party agent harnesses like OpenClaw; OpenAI hasn't—and Opus 4.6 token usage is down significantly while GPT-5.4's has surged. Plus: why the technical/non-technical split is the wrong way to think about AI adoption, who counts as an “author” when AI does the drafting, and a two-step design workflow from Every's team.

“How We Run a 25-person Company on Four AI Agents” by *Katie Parrott/Source Code*: Every runs six products, a media company, and a consultancy—and until recently, COO **Brandon Gell** was the router keeping all of it coordinated. Now four custom Notion agents handle prioritization, meeting-to-task conversion, OKR planning, and daily growth reporting. Read this for the full breakdown of each agent, and copy-paste prompts to build your own. (This piece was based on a camp sponsored by Notion.)

“Every Is Half Agent Now” by *Laura Entis/Context Window*: Every gave each employee a **Plus One**—a dedicated AI agent—and we're writing the etiquette for them as we go. **Brandon Gell** and **Willie Williams** join **Dan Shipper** to share what they've learned: Agents earn trust by executing tasks publicly, and everyone is a manager now whether they've had direct reports or not. Plus: Anthropic has built a powerful new model it's not releasing publicly; 70 percent of Every staff use gendered pronouns for their agents; and a prompt for when your agent won't stop talking. 🎧 🖥️ Listen on [Spotify](#) or [Apple Podcasts](#), or watch on [X](#) or [YouTube](#).

“The Market for Making AI Better” by *Alex Duffy/Thesis*: Reddit, Shutterstock, and News Corp are making hundreds of millions licensing data to AI labs, with contracts growing 20 percent annually. **Alex Duffy** argues that that undersells it: A 4-billion-parameter model recently beat one 60 times its size by training on the right financial data. Read this to understand what makes your company's proprietary data valuable, and whether to license it, train on it yourself, or both.

Log on

We host [camps and workshops](#) on topics like [compound engineering](#) and [writing with AI](#) to share the knowledge we've acquired from training teams at companies like the [New York Times](#) and leading hedge funds, and by learning and playing with AI every day ourselves.

Upcoming course

- [Claude Code for Absolute Beginners](#) (April 14): This beginner-friendly, live workshop led by **Mike Taylor** (head of tech consulting at Every) is designed to get you from zero to a working project with Claude Code. [Learn more](#).

From Every Studio

Sparkle is getting a full makeover

The [Sparkle](#) team has been working on a ground-up user interface redesign—new animations, new onboarding, new everything. General manager [Yash Poojary](#) says it doesn't even feel like the same app. The new version is already available to download [from Sparkle's website](#). Tune in next week for the full rollout.

Monologue Notes is live

[Monologue](#) now saves and organizes your recordings as browsable notes. General manager [Naveen Naidu](#) has been using it to capture everything from team calls to solo idea sessions, then pulling those notes into other tools via Monologue's CLI. The summaries are designed for builder workflows where you want to revisit what you were thinking, not just what you agreed to do. [Update](#) to the latest version to try it.

Spiral is experimenting with agent-to-agent workflows

Two days after the release of Anthropic's new [managed agents](#), [Marcus Moretti](#), general manager of [Spiral](#), has set them up to power Spiral's API. The setup lets an external agent (rather than a human) hand off a writing task to Spiral, where the two agents interview each other behind the scenes before producing a draft with no human input required. Marcus built a new API endpoint for this flow and added an API label in Spiral's UI so users can distinguish between agent-generated and human-initiated conversations. The API also now supports attachments and smarter default selections for workspace and style. Conversations via API show up in your Spiral chat history with an "API" label, so you can pick up where the agent left off.

Alignment

The wrong fight. I don't know what's in the water in Utah, but whatever it is, I want more of it, because the state is leading the country on using AI in healthcare.

Legion Health, a Y Combinator-backed San Francisco startup, has been cleared to use AI in Utah to renew a handful of psychiatric prescriptions, including Prozac and Zoloft, for patients

who are already stable and on an established treatment plan. It's the second AI healthcare pilot approved there, and it's replacing the barrages of emails from patients who are stable on the same dose, contacting their clinicians who are already buried in administrative work, who have to produce a piece of paper that says yes, same drug, same dose, carry on. This is often done outside of working hours, and without any reimbursement.

To ensure the pilot is safe, the first 250 AI renewals are reviewed by a physician before anything reaches a pharmacy, and the AI has to agree with that physician more than 98 percent of the time before it can proceed independently. The next 1,000 renewals are then reviewed, with an even higher threshold of 99 percent before the oversight shifts to randomized monthly testing, with Legion filing monthly reports on accuracy and any adverse outcomes throughout.

Yet both the tech coverage and members of the medical establishment have deemed it too risky. The criticism splits into two camps: prescribing error, and the app's insufficiency to improve access to the patients who need care most. On prescribing error, the hard clinical judgment has already been made by a human; what the AI is doing is confirming that nothing has changed, which it has to get right 98 percent of the time before it's allowed to proceed unsupervised.

On access, it's true that you have to already be in treatment to use this service, but if a psychiatrist in rural Utah who typically spends part of their day processing renewal emails for stable patients no longer needs to do so, they have more time for the patients who need them.

Most of Utah's counties are designated mental health provider shortage areas, leaving around 500,000 residents without adequate psychiatric care.

Physician risk-aversion is one of medicine's great virtues in the right context, but renewing a stable prescription is not that context, and dressing up administrative inertia as a patient safety concern doesn't make it one.—Ashwin Sharma

That's all for this week! Be sure to follow Every on X at @every and on LinkedIn.

We build AI tools for readers like you. Write brilliantly with Spiral. Organize files automatically with Sparkle. Deliver yourself from email with Cora. Dictate effortlessly with Monologue. Work on documents with AI agents using Proof.

For sponsorship opportunities, reach out to

sponsorships@every.to

.

Good morning, tech enthusiasts. Snap is taking another shot at smart glasses — this time with Qualcomm chips, on-device AI, and real pressure to deliver.

The reboot follows the abrupt exit of its top Specs exec amid a reported clash with CEO Evan Spiegel, as the company races to beat Meta to the face-worn AI interface. But Meta's Ray-Ban glasses are already gaining traction; Snap's Spectacles are still trying to ship.

In today's tech rundown:

- - Snap's Spectacles get a Qualcomm engine
- Tesla may build a low-priced SUV after all
- Meta bans the ads being used to sue it
- One therapy wipes out 3 autoimmune diseases
- Quick hits on other tech news

LATEST DEVELOPMENTS

SNAP

Snap's Spectacles get a Qualcomm engine

Image source: Snap

The Rundown: Snap is finally inching toward launching its long-delayed AI-powered Spectacles, striking a multi-year Qualcomm chip deal to bring its next-gen AR glasses to consumers later this year after a recent executive shake-up at its Specs unit.

The details:

- - The new Spectacles will run on Qualcomm's Snapdragon XR chips, enabling on-device AI, advanced graphics, and multiuser digital experiences.
- Snap spun out Specs as a separate subsidiary earlier this year to focus on the glasses business after years of fits and starts with the product line.
- In February, Snap abruptly parted ways with Scott Myers, its senior VP of Specs, following a reported clash with CEO Evan Spiegel.

Why it matters: Snap's decade-old Specs unit is under pressure to finally turn Spectacles into a real consumer platform just as Meta, Apple, and others race to dominate the smart glasses market. If this launch fizzles like past attempts, Snap risks ceding the next hardware frontier to rivals with deeper pockets and tighter ecosystems.

TESLA



Tesla may build a low-priced SUV after all

Image source: Reve / The Rundown

The Rundown: Tesla is developing a compact electric SUV designed to undercut its own Model 3 on price, Reuters [reports](#) — two years after CEO Elon Musk scrapped the \$25K "Model 2" and called building cars for human drivers "pointless."

The details:

- The vehicle would measure about 14 ft., making it significantly shorter than the Model Y's 15.7 ft., and would be an entirely new design.
- To hit a lower price point, Tesla plans to use a smaller battery pack and a single electric motor, trading range against the Model Y's 306-to-327-mile rating.
- Sources said production would be based at Tesla's Shanghai factory, with one source adding that Tesla aims to expand manufacturing to the U.S. and Europe.
- Pricing would land substantially below the entry-level Model 3, which starts at \$34K in China and \$37K in the U.S.

Why it matters: Tesla's sales have taken a hit as Chinese EVs flood the sub-\$30K segment, a price point the company has never actually reached. A compact SUV would be its most direct answer yet to that pressure, though with no formal approval and only early supplier conversations underway, the timeline remains unclear.

META



Meta bans the ads being used to sue it

*Image source: Getty /
Reve*

The Rundown: Meta just removed a wave of Facebook and Instagram ads placed by plaintiffs' law firms recruiting clients for social media addiction lawsuits, as litigation against the company continues to mount.

The details:

- - Meta has pulled campaigns from major national firms targeting teens and parents to join social media addiction lawsuits.
- Meta said it wouldn't "allow trial lawyers to profit from our platforms while simultaneously claiming they are harmful."
- An LA jury recently awarded \$6M against Meta and Google for a woman's depression; a New Mexico jury fined Meta

\$375M over child safety failures.

- More than 3,300 addiction-related lawsuits are pending in California state courts, with another 2,400 federal cases centralized there.

Why it matters: The same ad-targeting engine that helped Meta capture attention is now being used to recruit people to sue them. With billions in potential liability on the line and thousands of cases pending, the move shows how fiercely Meta intends to fight a legal battle that could reshape platform accountability for teen mental health.

BIOTECH



One therapy wipes out 3 autoimmune diseases

Image source:

Ideogram / The

Rundown

*(rendering of a
CAR-T cell)*

The Rundown: For the first time, a single round of experimental CAR-T-cell therapy **put** all three of a patient's severe, treatment-resistant autoimmune diseases into lasting remission — a result doctors say they've never seen before.

The details:



The woman was managing three debilitating autoimmune conditions until a single infusion of T cells effectively rebooted her immune system.

- Doctors hacked a blood-cancer treatment, reprogramming her T cells to hunt down CD19-tagged B cells, the antibody factories gone rogue in her system.
- Within weeks, her blood counts normalized as a fresh population of mostly naïve B cells repopulated her system.
- Fourteen months on, she remains off all medications for the three conditions, with no reported side effects from the therapy itself.

Why it matters: CAR-T-cell therapy has already **transformed** blood cancer treatment — but repurposing it for autoimmune disease is a newer, bolder bet. A single case isn't a cure, and larger trials are needed. Still, sustained triple remission without ongoing medication is the kind of outcome researchers rarely dare to predict.

QUICK HITS



Everything else in tech today

SpaceX **recorded** a loss of nearly \$5B in 2025 despite generating more than \$18.5B in revenue, according to a report by The Information.

Disney's new CEO, Josh D'Amaro, **plans** to cut up to 1K jobs, with the company's recently consolidated marketing department expected to bear the brunt of the layoffs.

Almost half of the U.S. data centers **planned** to open in 2026 are likely to be delayed or canceled due to power grid limits, equipment shortages, and local opposition.

Apple's first foldable iPhone **remains** on track to be unveiled in September 2026 alongside the iPhone 18 Pro lineup, despite earlier rumors of delay, Bloomberg reports.

Spotify is **adding** new global settings so you can turn off all videos, including music videos, podcast videos, and Canvas loops, and keep Spotify audio-only if you want.

Instagram is **rolling out** a long-requested feature that lets users edit their comments shortly after posting, so they no longer need to delete and rewrite them.

Volkswagen will stop **building** its ID.4 electric SUV at its Tennessee plant and shift the factory to producing the higher-volume, gasoline-powered Atlas SUV instead.

NASA **traced** the helium leak in Orion's propulsion system to faulty valves and is now planning a hardware redesign to prevent similar issues on future lunar missions.

The EU has **hit** Google, Apple, and Meta with more than \$7B in antitrust

and digital regulation fines since 2024, triggering a clash with the U.S. government.

Startup Radify Metals

developed plasma reactors that can refine rare-earth metals, potentially undercutting China's dominance over the rare-earth supply chain.

Greece **plans** to ban social media access for children under 15 starting January 1, 2027, as part of new legislation aimed at protecting young people's mental health.

Finland's long-planned Onkalo facility is on the verge of **becoming** the world's first operational deep geological repository for spent nuclear fuel.

COMMUNITY



Highlights: News, Guides & Events

- Read our last AI newsletter: [Perplexity's agent pivot is on the money](#)
- Read our last Tech newsletter: [This startup wants to hack the](#)

[night](#)

[sky](#)

- Read
our last
Robotics
newslett
er:

[Clone's](#)

[\\$20k](#)

[syntheti](#)

[c](#)

[human](#)

- Today's
AI tool
guide:

[Automa](#)

[te your](#)

[busines](#)

[s with](#)

[custom](#)

[Notion](#)

[Agents](#)

- Watch
our last
live
worksho
p: [The](#)
[State of](#)
[AI](#)

Present

ation

Tools in

2026

**That's it
for
today's
tech
rundown!**

We'd love to
hear your
feedback on
today's
newsletter so
we can
continue to
improve The
Rundown
experience for
you.

Good morning, AI enthusiasts. Perplexity stopped trying to simply out-Google Google and started going after Mint, TurboTax, and every personal finance app on your phone instead.

The company's Computer launch in February now doubles as a personal finance hub and tax tool, changing Perplexity's trajectory from search to something much bigger — with a 50% monthly revenue jump to \$450M backing up every agentic move.

In today's AI rundown:

- - Perplexity plugs its AI agent into bank accounts
- Jassy's \$200B Amazon AI spend now has receipts
- Automate your business with custom Notion Agents
- Oxford AI catches heart failure five years early
- 4 new AI tools, community workflows, and more

LATEST DEVELOPMENTS

PERPLEXITY



[Perplexity plugs its AI agent into bank accounts](#)

Image source: Perplexity

The Rundown: Perplexity just [rolled out](#) a new Plaid integration that lets users connect bank accounts, credit cards, and loans directly to its Computer agent, turning it into a full personal finance hub.

The details:

- - Plaid's 12K+ bank network feeds into Computer, with users able to pull in checking, credit, loan, and brokerage data for a read-only view of their money.
- The agentic system can then build customized tools like budgets, net worth trackers, debt payoff plans, and retirement dashboards via simple text prompts.
- The move comes on the heels of Perplexity's U.S tax [integration](#) that autonomously fills out IRS forms and reviews professional-prepared returns.
- Perplexity Computer [launched](#) in late February, with the agentic pivot helping push Perplexity's ARR past \$450M in March, a 50% jump in a single month.

Why it matters: Perplexity built its name trying to out-Google Google, but it's Computer has completely changed the trajectory. With smart connectors and a powerful AI agent, the company is suddenly competing with Mint, TurboTax, and every other app area it ends up integrating — not just search.

TOGETHER WITH SERPAPI



[Access Google Search, Maps, and more live data via API](#)

The Rundown: Need real-time public web data without scraper headaches? SerpApi turns search results into structured, ready-to-use data with built-in scale and block resilience — powering AI apps, product research, price tracking, SEO insights, and more without the maintenance.

With SerpApi, you get:

- Real-time, reliable results from any search engine
- ZeroTrace Mode to keep your searches confidential

- U.S. Legal Shield for scraper legal protection

[Try SerpApi playground for free today](#) —

Rundown readers can get 50% off for 3 months via chat or email.

AMAZON

[Jassy's \\$200B Amazon AI spend now has receipts](#)

Image source: Amazon

The Rundown: Amazon CEO Andy Jassy [shared](#) his annual shareholder letter with the company's first-ever AI revenue figures and a defense of the \$200B planned capex, dismissing bubble talk and floating the idea of selling Trainium chips to outside buyers.

The details:

- Amazon's \$200B AI spending rattled investors this year, with Jassy's letter firing back with first-ever revenue figures and locked-in customer demand.

- AWS's AI arm crossed \$15B in annualized revenue, a number Amazon had never disclosed — and 260x where AWS itself stood at the same point.
- The custom Trainium, Graviton, and Nitro chips crossed \$20B in yearly revenue, and Amazon may sell “racks of them to third parties in the future.”
- Two unnamed AWS customers asked to buy the company's entire Graviton chip supply for 2026, with Amazon declining to protect other clients' access.

Why it matters: If you only tracked models as a barometer for the AI race, Amazon might look like it's behind — but the \$20B chip numbers tell a different story. Nvidia has dominated AI compute, but the supply side of the boom is finally getting real competition at exactly the moment demand has never been higher.

Automate your business with custom Notion Agents

The Rundown: In this guide, you will learn how to build Notion Custom Agents that run your company's recurring work on a schedule, automating inbound leads, campaigns, accounting, or any recurring job in your business without adding complexity.

Step-by-step:

1.
In [Notion](#), create two databases — Tasks (Name, Source, Priority, Status, Assigned To) for every to-do, and Reports, so that agents can log their work
2. Open Notion AI, click +Create custom agent, and prompt: "Create a Weekly Planner agent that reads my last 7 days of

emails every Monday morning, adds action items into the Tasks database, and writes a one-paragraph summary of the run into the Reports database”

3. Notion drafts the agent with a name, trigger, instructions, and data sources. Review, connect Gmail, confirm the schedule, and save
4. Run it on demand to test. Tighten the instructions if the output is off, then leave it running. Notion shows the next scheduled run, and the work is off your plate

Pro tip: Clone the pattern for every recurring job. Same two databases, different agent, with every run traceable in one place.

PRESENTED BY UNFRAME



Enterprise AI

ROI: 2026

benchmarks

The Rundown: While enterprise AI is widely deployed and already delivering measurable gains, many organizations struggle with translating those gains into business outcomes.

Unframe surveyed 255 enterprise leaders to reveal current trends in ROI, adoption, and scaling of AI.

In this report, you'll find:

- 4 in 5 enterprises report productivity gains
- ROI drops 25% in environments with 6+ tools
- Half of potential value is lost between insight and action

[Get the report.](#)

AI MEDICAL RESEARCH



[Oxford AI](#)
[catches heart](#)
[failure five years](#)
[early](#)

Image

source:

Lovart /

The

Rundown

The Rundown:

Researchers at the
University of Oxford
[introduced](#) an AI system
that picks up invisible
changes in heart fat from
routine CT scans, flagging
patients at high risk of
heart failure up to five
years out — with 86%
accuracy across 72K
patients.

The details:

- Fat around the heart shifts texture when the muscle beneath is inflamed, with the AI reading the patterns invisible to doctors on any current scan.
- In the highest-risk bucket, 1 in 4 patients ended up with heart failure within five years — a 20x gap versus those the AI flagged as safe.
- Oxford is already working with regulators to bring the tool to National Health Service hospitals, and plans to extend

it to all chest CT
scans within
months.

Why it matters: Heart failure's biggest problem isn't treatment, it's timing. Doctors usually can't act until damage has set in, so an 86%-accurate early warning system built into scans patients are already getting could shift the equation of a serious condition from reaction to prevention for better diagnosis and outcomes.

QUICK HITS

 [Trending](#)
[AI Tools](#)

•



[Claude](#)

Cowork

-

Anthropi

c's

desktop

agent,

now

generall

y

available

- 

Perplexi

ty

Comput

er - New

agentic

connecti

ons for

finances

, taxes

- 

Muse

Spark -

Meta's

multimo

dal

reasonin

g AI with

multi-ag

ent

mode

- 🐱

[Meow](#) -

Infra that

lets

agents

open

bank

account

s, issue

cards,

and

more



[Everything
else in
AI today](#)

Spacelift

Intelligence

just launched,

an AI

infrastructure

suite that helps

platform teams

ship infra as

fast as
developers
code. [Start for
free](#).*

OpenAI has
[built](#) a model
with advanced
cybersecurity
skills similar to
Anthropic's
Mythos, with
Axios reporting
the company
plans to
release it to a
"small set of
partners".

xAI is
[undergoing](#) a
reorg of its
engineering
division, with
CFO Anthony
Armstrong
[leaving](#) the
company as
SpaceX execs
are installed
ahead of the
company's
IPO.

OpenAI
[launched](#) a

\$100/month

Pro tier with 5x

more Codex

usage than

Plus, designed

for heavy

agentic coding,

coming amid

anger over

Claude usage

limits.

Florida's

attorney

general

[opened](#) a

probe into OAI

with

subpoenas

incoming,

citing

allegations that

ChatGPT

helped plan a

campus

shooting at

Florida State

University.

**Sponsored*

Listing

COMMUNIT

Y



[Comm](#)
[unity](#)

[AI](#)
[workfl](#)
[ows](#)

Every
newslette
r, we
showcas
e how a
reader is
using AI
to work
smarter,
save
time, or
make life
easier.

Today's
workflow
comes
from
reader

A.M in

New

Zealand:

"I had an
interview
for a role
that was
a leap up
from

where I
was
currently
at. I

loaded
the job
ad into
ChatGPT
and

asked it
to run me
through
interview
questions
that I

voice-ans
wered.

Chat
reviewed
my

answers,
and I

repeated
this
process

until I
scored
highly.

I then
made it
to the
2nd
round of
interview
s with the
board. I
research
ed all 5
board
members
and then
had Chat
run me
through
another
list of
questions
I could
expect,
and then
repeated
the
process
of
reviewing
my
answers

until I
scored
highly.
Bonus
points for
Chat
helping
me
answer
questions
in a way
that
would
appeal to
the
specific
board
members
' roles
and
interests.
Fingers
crossed
for a
response
any day
now!"
How do
you use
AI? Tell
us [here](#).

[Aman Sanger, Cursor: “After switching to turbopuffer, we're saving an order of magnitude in costs” \(Sponsor\)](#)

turbopuffer is the [serverless search engine for AI](#). It's built on object storage with smart caching - so it's fast, cheap, and extremely scalable.

✓ Used by Cursor, Notion, Linear, Cognition, Atlassian, Ramp, Granola, Legora...

✓ Fast → ~20ms p90 latency on 10M documents

✓ Truly serverless → [Run your first query in ~4 minutes](#)

✓ Up to 95% cheaper than traditional vector/text search databases

TLDR readers get their first month free. [Sign up here](#).

[ChatGPT \\$100 Pro Plan Introduced \(2 minute read\)](#)

OpenAI introduced a new \$100/month ChatGPT Pro tier aimed at power users, expanding its pricing lineup between the \$20 Plus and \$200 Pro plans. The company confirmed the \$200 plan still exists despite not being listed on the pricing page.

[Making Claude Cowork ready for enterprise \(3 minute read\)](#)

Claude Cowork is now enterprise-ready, enhancing work with organization controls like role-based access, group spend limits, and expanded observability features. Admins can manage adoption with detailed usage analytics and integrate tools like Zoom for seamless workflows. Companies like Zapier and Airtree have already leveraged these features for improved project management and operational efficiency.

Alibaba Claims Viral Happy Horse AI Model in Latest Breakthrough **(3 minute read)**

The Happy Horse video AI model that sent ripples across China's AI industry this week was created by Alibaba. The model hit the top spot on the text-to-video leaderboard of Artificial Analysis this week. Alibaba plans to provide API access to the model in the near future.

Perplexity launches Personal Finance powered by Plaid (2 minute read)

Perplexity has expanded its financial services with Plaid integration, transitioning from portfolio tracking to a comprehensive finance dashboard. Users can link checking, savings, credit card, and loan accounts to analyze spending, track liabilities, and calculate net worth. Targeted at AI-savvy consumers, it combines data insights with Plaid's infrastructure, transforming from a portfolio add-on to a full personal finance hub without executing trades.

Deep Dives & Analysis

[Agentic Infrastructure on Vercel \(8 minute read\)](#)

Vercel argued that coding agents were reshaping how software gets built and deployed, with agent-initiated deployments rising sharply and accounting for more than 30% of weekly deployments. The post framed this shift as a need for new infrastructure designed for agents to deploy software, run AI systems, and increasingly operate infrastructure autonomously.

[CoreWeave Takes As Much Financial Engineering As It Does Datacenter Design \(4 minute read\)](#)

CoreWeave's revenue backlog is up to \$87.8 billion, with Meta represents 40.1% of that backlog. OpenAI's deals represent 25.5% of extended backlog. CoreWeave only has two dozen named customers, but it likely has more. The company achieved \$5.13 billion in sales in 2025, up 2.7x from 2024, and posted a net loss of \$1.17 billion. It recently held a private offering of \$1.75 billion in notes to cover its data center expansion costs.

[Research-Driven Agents: What Happens When Your Agent Reads Before It Codes \(16 minute read\)](#)

Coding agents generate better optimizations when they read papers and study competing projects before touching code. Coding agents working from code alone generate shallow hypotheses. Adding a research phase improves results.



Engineering & Research

[IronClaw: the secure agent harness for always-on AI \(Sponsor\)](#)

You're one API key away from an agent that codes, commits, and pulls from GitHub and Slack autonomously. What's stopped you is trusting it with credentials. [IronClaw](#) isolates every tool in a Wasm sandbox. Credentials never reach the model.

Open-source, built in rust, deployable on NEAR AI Cloud. Verifiably yours.

[Anthropic launches advisor tool for Claude Platform API users \(1 minute read\)](#)

The new advisor tool on Anthropic's Claude Platform API provides developers with the ability to use Opus as an advisor alongside Sonnet or Haiku as executors. This allows agents to access advanced reasoning capabilities while maintaining operational costs at the level of the more efficient executor models. The advisor tool is now publicly available through a configuration in the Messages API request.

Multimodal Embedding & Reranker Models with Sentence Transformers (8 minute read)

The v5.4 update of Sentence Transformers introduces multimodal embedding and reranker models that enable encoding and comparing texts, images, audio, and videos within a shared embedding space for tasks like cross-modal search and retrieval-augmented generation.

These advanced models map different modalities into a consistent embedding space and assess relevance across mixed-modality pairs, though lower cross-modal similarity scores may occur due to modality clustering. Notably, multimodal rerankers provide high-quality scoring by evaluating individual pairs for more accurate retrieval results.

Process Driven Image Generation (22 minute read)

Meta has introduced a multi-step framework for image synthesis that alternates between textual reasoning and visual updates. This enables models to iteratively plan, draft, reflect, and refine outputs.

Efficient Rollout Scaling for Diffusion RL (4 minute read)

NVIDIA's Sol-RL introduces a two-stage framework that separates exploration and training, using FP4 rollouts to generate large candidate sets and BF16 for selective policy updates. This approach reduced compute costs while improving alignment and accelerating convergence in diffusion model post-training.

[Dario Says Continual Learning Is Solved. Is It? \(5 minute read\)](#)

Continual learning research is aimed at breaking past the feasible horizon of current techniques. If models can't learn new things while performing tasks, they will struggle when the task horizon grows very long. Context management and other engineering improvements can push task horizons to weeks, or even months, but very long-horizon tasks remain out of reach. Human-level continual learning may be 'solved', but the revolution is not yet complete.

[Introducing KellyBench \(2 minute read\)](#)

KellyBench evaluates sequential decision-making in sports betting by simulating the 2023-24 English Premier League season. Models like Claude Opus 4.6 and GPT-5.4 struggled, with none achieving positive returns, highlighting their limitations in adapting to long-term strategies. This underscores a need for complex environments enabling agents to learn from experience under uncertainty.

[The Gemini app can now generate interactive simulations and models \(1 minute read\)](#)

The Gemini app now supports generating interactive visualizations, like 3D models and simulations, directly in chat.

[Google's PaperOrchestra AI Converts Lab Notes Into Publication-Ready Research Papers \(2 minute read\)](#)

Google Cloud AI's PaperOrchestra converts disorganized lab notes into submission-ready academic papers using five specialized agents.

OpenAI Projects Steep Advertising Growth, Targeting \$100 Billion by 2030 (1 minute read)

OpenAI expects its advertising business to generate \$2.5 billion in revenue this year.